

A Recommendation System for Spotify

Luigi Barba, Lorenzo Colantoni, Francesco Guglielmino,
Jacopo Masini e Gianluca Surico

28/07/2020

Introduction

A recommendation system is a content filtering software that creates personalized recommendations specific to the user in order to help him in his choices.

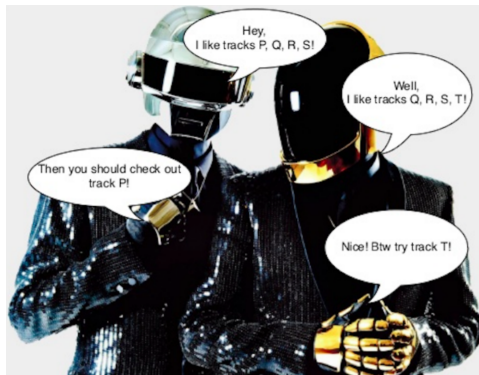


Introduction

Approaches

In general, recommendation systems are based on two different strategies:

- ▶ Content-based approach
- ▶ Collaborative Filtering (CF)



Introduction

Kinds of inputs

The recommendation systems are based on different types of inputs:

- ▶ the easiest to manage is explicit feedback, which includes explicit input from users about their interest in products.
- ▶ the easiest to get is implicit feedback, which indirectly reflects the user's opinion.

Main goal

The goal of our work is to recommend new artists to listen to users taken as input, using collaborative filtering with implicit feedback data. This consists in analyzing the relationship between user and artist, working on implicit signals as artist's track play counts.

Data collection

We collected data for 12 Spotify accounts, in particular the list of favorite tracks for each user. We have merged the data into an $m \times n$ matrix. The m rows represent the users, the n columns are the unique artists of all users' favorite tracks.

Model Overview (1) - The input data

We define the matrix $R = (r_{ui}) \in \mathbb{R}^{m \times n}$ as the user-item interaction matrix.

Each entry r_{ui} specifies the interaction of item i by user u and can be thought of as a single observation. Here's an example of R :

	Item 1	Item 2	Item 3
User 1	0	5	4
User 2	1	2	6
User 3	0	0	7

Underlying implicit feedback model hypothesis: r_{ui} is definitely a measure of appreciation, but it is rather "raw" and needs to be processed further!

Model Overview (2) - Building a confidence metric

For this reason we are interested in formalizing a measure of confidence in how the user really appreciate the item he interact with:

$$c_{ui} = 1 + \alpha \log(1 + r_{ui}) \quad (1)$$

Remark: the confidence function is not univocal, but can be chosen at discretion. The better it reflects the nature of the data, the better it improves prediction accuracy.

Multiple confidence functions (along with their parameters) could even be objects of cross-validation!

Model Overview (3) - Finding latent factor matrices

In Matrix Factorization models, the goal is to find a vector $x_u \in \mathbb{R}^f$ for each user u , and a vector $y_i \in \mathbb{R}^f$ for each item i that will factor user preferences.

In other words, preferences are assumed to be the inner products of two vectors known as user-factors and the item-factors, respectively:

$$p_{ui} = x_u^T y_i \quad (2)$$

where p_{ui} is a binary variable:

$$p_{ui} = \begin{cases} 0 & \text{if } r_{ui} = 0 \\ 1 & \text{if } r_{ui} > 0 \end{cases} \quad (3)$$

Model Overview (4) - The Optimization problem

Based on the what we have seen, we can finally define our Loss Function:

$$\min_{x_*, y_*} \sum_{u,i} c_{ui} (p_{ui} - x_u^T y_i)^2 + \lambda (\sum_u \|x_u\|^2 + \sum_i \|y_i\|^2) \quad (4)$$

The first term is the model intrinsic loss, the second one is the penalty assigned to the latent factors in order to prevent overfitting. λ , along with α (contained in c_{ui}), are two parameters to be tuned.

Model Overview (5) - Additional optimization I

This cost function contains $m \times n$ terms, where m is the number of users and n is the number of items. For typical datasets (even if it is not this project's case) this product can easily reach a few billion. What to do?

Model Overview (6) - Additional optimization II

Solution: by exploiting the fact that when either the user-factors or the item-factors are fixed the cost function becomes quadratic (so its global minimum can be readily computed), we could consider an alternating-least-squares optimization process, where we alternate between re-computing user-factors and item-factors, and each step is guaranteed to lower the value of the cost function.

- ▶ Step 1: Compute Y
- ▶ Step 2: Fix Y from Step 1 and compute X
- ▶ Step 3: Fix X from Step 2 and compute Y
- ▶ And so on, until it stabilizes

Model Overview (7) - Additional optimization III

This model's version of the alternating-least-squares optimization is peculiar since it requires to take into account the confidence levels we defined before. In fact, X and Y are computed as follows:

$$x_u = (Y^T C^u Y + \lambda I)^{-1} Y^T C^u p(u) \quad (5)$$

$$y_i = (X^T C^i X + \lambda I)^{-1} X^T C^i p(i) \quad (6)$$

where C^u is a the diagonal $n \times n$ matrix with $C_{ii}^u = c_{ui}$, and $p(u)$ is a vector $\in \mathbb{R}^n$ that contains all the preferences by u (the p_{ui} values)

The implementation

Our idea is to split the dataset into training and test set and then artificially split the test set into history and future data in order to quantify the performance of the model.

In R this will be done using the `createDataPartition` of the library `caret`.

The design matrix

The design matrix X is a user-item matrix in which the r_{ui} element is the number of saved songs made by artist i saved by user u .

The feedback is implicit (no explicit rating).

In R, this will be saved in the 'sparseMatrix' format mainly because of the sparsity of X .

The Tuning

The tuning of the parameters will be done using a 4-fold cross-validation scheme.

There are 3 tuning parameters: the number of latent factors, the α confidence parameter and the regularization parameter λ .

The fitting of the model is made possible by using the *model\$fit_transform* function of the *rsparse* package after setting up the model using *model = WRMF\$new()*.

The Evaluation Metric

The validation scheme is based on a precision measure called $MAP@k$: it is the average over the test set of the average precision for each user.

$$AP@k = \sum_{k=1}^n \frac{P(k) * rel(k)}{I(u)} \quad (7)$$

In R, this is done using the `ap_k` function of the `rsparse` package, which returns the average precision of the predictions for each user in the test set: the $MAP@k$ will be the mean of this vector


The recommendations

After tuning the parameters, the recommended items will be displayed after using the *model\$prediction* function, which - on a mathematical level - just performs another round of the ALS algorithm.


Cross Validation and Parameters Tuning

Before starting analyzing the results we got, let's remark briefly an important concept on which our work is based on:


- ▶ the Cross Validation we applied isn't a "pure" one based on the usually loss function \longrightarrow it's based on the **MAP@k** metric evaluation that is independent from α parameter



Boxplot of Average CV-performance by rank (tuning λ & α)



Average CV-performance by rank (tuning λ & α)



Model Recommendation

Lorenzo		Jacopo		Gianluca	
Score	Recommendation	Score	Recommendation	Score	Recommendation
0.8517400	Arctic Monkeys	0.5618116	Bruce Springsteen	0.7689953	Coez
0.7965103	Bruno Mars	0.5284429	Simon & Garfunkel	0.7043944	Simon & Garfunkel
0.7849014	Black Eyed Peas	0.4720040	Coez	0.6490270	J-AX
0.7609935	Billie Eilish	0.4713720	J-AX	0.6147510	Bob Dylan
0.7443408	Sia	0.4706330	The Clash	0.5899966	Linkin Park
0.7443408	Michael Jackson	0.4620708	U2	0.5757663	U2
0.7337686	Mannarino	0.4590332	The Beatles	0.5708236	Jefferson Airplane
0.7205724	Eminem	0.4552083	Linkin Park	0.5637806	The Doors
0.7071880	Rino Gaetano	0.4484908	Ludovico Einaudi	0.5576449	The Smashing Pumpkins
0.6844891	Nitro	0.4376672	alt-J	0.5376591	Jarabe De Palo

Similarity between artists

- ▶ recommended artists
- ▶ similar artists

Look for similarities between the items of our artists' basin.

After implementing our recommendation system we tried to exploit the information coming from the "latent" item matrix in order to suggest to the users artists similar to the one already recommended.

This was done using the `rsparse_model$components`, computing cosine similarities between the columns.

Similarity between artists

Cosine similarity:

$$sim_{ab} = \frac{\sum_i a_i b_i}{\sqrt{\sum_i a_i^2} \sqrt{\sum_i b_i^2}} \quad (8)$$

where a , b represent the vectors in a latent factor space for two artists.

Similarity between artists

Latent factor			
The Doors	David Bowie	Jefferson Airplane	Bruce Springsteen
0.20757844	0.20052335	0.16216339	0.167778510
0.18957669	0.17499336	0.18517567	0.170741923
0.20070742	0.27311633	0.10546285	0.100863064
-0.17000637	-0.09837597	-0.11116437	-0.088528736
-0.13047577	0.14050596	-0.06114843	-0.059720767
0.05400830	-0.13187250	-0.13761145	-0.119644873
-0.01616389	-0.07034186	-0.03062981	-0.003315523

Similarity between artists

Similar artists		
The Doors	The xx	Iron Maiden
David Bowie	Beck	Lou Reed
Jefferson Airplane	Talking Heads	alt-J
Bruce Springsteen	Little Simz	The Rolling Stones
Bracket	Fleet Foxes	Phil Collins
Lucio Battisti	Massive Attack	Caravan

Similarity between artists

